KI Service Zentrum by Hasso-Plattner-Institut

HPI

Gefördert durch:

Bundesministerium für Forschung, Technologie und Raumfahrt

Effizientes Abfragen von Informationen II

Lokales Retrieval Augmented-Generation System

Felix Boelter
Hanno Müller

Design IT.
Create Knowledge.

www.hpi.de

Vier AI Service-Zentren in Deutschland

KI Service Zentrum by Hasso-Plattner-Institut HPI

https://hpi.de/ki-servicezentrum/

Überblick

KI-Servicezentren

Ziel:
Reduktion von Barrieren bei der Implementierung von KI-Anwendungen in Gesellschaft und Wirtschaft.

Gefördert durch:
Bundesministerium für Forschung, Technologie und Raumfahrt

KI-Servicezentrum Berlin Brandenburg

KI Service Zentrum by Hasso-Plattner-Institut HPI

Struktur

Gefördert durch:

Bundesministerium für Forschung, Technologie und Raumfahrt

KI Service Zentrum by Hasso-Plattner-Institut

HPI

Forschung

- **KI-Methodenforschung**
- **KI-Betriebsforschung**

PubMedCLIP: How Much Does CLIP Benefit Visual Question Answering in the Medical Domain?

Sedigheh Eslami, Christoph Meinel, Gerard de Melo
Hasso Plattner Institute / University of Potsdam
{sedigheh.eslami, christoph.meinel, gerard.demelo}@hpi.de

Abstract

Contrastive Language–Image Pre-training (CLIP) has shown remarkable success in learning with cross-modal supervision from extensive amounts of image–text pairs collected online. Thus far, the effectiveness of CLIP has been investigated primarily in general-domain multimodal problems. In this work, we evaluate the effectiveness of CLIP for the task of Medical Visual Question Answering (MedVQA). We present PubMedCLIP, a fine-tuned version of CLIP for the medical domain based on PubMed articles. Our experiments conducted on two MedVQA benchmark datasets illustrate that PubMedCLIP achieves superior results improving the overall accuracy up to 3% in comparison to the state-of-the-art Model-Agnostic Meta-Learning (MAML) networks pre-trained only on visual data. The PubMedCLIP model with different back-ends, the source code for pre-training them and reproducing our MedVQA pipeline is publicly available at https://github.com/sarahESL/PubMedCLIP.

1 Introduction

Medical visual question answering (MedVQA) seeks answers to natural language questions about a given medical image. The development of MedVQA has considerable potential to benefit healthcare systems, as it may aid clinicians in interpreting medical images and obtaining more accurate diagnoses by consulting a second opinion. Thus, it has become a very active area of research, with competitive benchmarks and yearly competitions (Abacha et al., 2021). Yet, visual question answering in the medical domain in particular remains non-trivial as we suffer from a general lack of large balanced training data, in part due to privacy concerns. To solve the multimodal task of MedVQA, a system must understand both medical images and textual questions and infer the associations between them sufficiently well to produce a correct answer (An

Findings of the Association for Computati
May 2-6, 2023 ©2023 Associ

Exploring Paracrawl for Document-level Neural Machine Translation

Yusser Al Ghussin[1,2], Jingyi Zhang[3] and Josef van Genabith[1,2]
[1]German Research Center for Artificial Intelligence (DFKI),
Saarland Informatics Campus, Saarbrucken, Germany
[2]Department of Language Science and Technology, Saarland University, Germany
[3]Hasso-Plattner-Institut (HPI), Potsdam, Germany
yusser.al_ghussin/Josef.Van_Genabith@dfki.de,Jingyi.Zhang@hpi.de

Abstract

Document-level neural machine translation (NMT) has outperformed sentence-level NMT on a number of datasets. However, document-level NMT is still not widely adopted in real-world translation systems mainly due to the lack of large-scale general-domain training data for document-level NMT. We examine the effectiveness of using Paracrawl for learning document-level translation. Paracrawl is a large-scale parallel corpus crawled from the Internet and contains data from various domains. The official Paracrawl corpus was released as parallel sentences (extracted from parallel webpages) and therefore previous works only used Paracrawl for learning sentence-level translation. In this work, we extract parallel paragraphs from Paracrawl parallel webpages using automatic sentence alignments and we use the extracted parallel paragraphs as parallel documents for training document-level translation models. We show that document-level NMT models trained with only parallel paragraphs from Paracrawl can be used to translate real documents from TED, News and Europarl, outperforming sentence-level NMT models. We also perform a targeted pronoun evaluation and show that document-level models trained with Paracrawl data can help context-aware pronoun translation. We release our data and code here[1].

1 Introduction

The Transformer translation model (Vaswani et al., 2017), which performs sentence-level translation based on attention networks, has achieved great success and significantly improved the state-of-the-art in machine translation. Compared to sentence-level translation, document-level translation (Xu et al., 2021; Bao et al., 2021; Jauregi Unanue et al., 2020; Ma et al., 2020; Maruf et al., 2019; Tu et al., 2018; Maruf and Haffari, 2018) performs translation at document-level and can potentially fur-

[1]https://github.com/Yusser96/Exploring-Paracrawl-for-Document-level-Neural-Machine-Translation

12

Proceedings of the 17th Conference of the European Chapter of
May 2-6, 2023 ©2023 Associati

Efficient Parallelization Layouts for Large-Scale Distributed Model Training

Johannes Hagemann
Aleph Alpha / Hasso Plattner Institute
johannes.hagemann@student.hpi.de

Samuel Weinbach
Aleph Alpha
samuel.weinbach@aleph-alpha.com

Konstantin Dobler
Hasso Plattner Institute
konstantin.dobler@hpi.de

Maximilian Schall
Hasso Plattner Institute
maximilian.schall@hpi.de

Gerard de Melo
Hasso Plattner Institute
gerard.demelo@hpi.de

Abstract

Efficiently training large language models requires parallelizing across hundreds of hardware accelerators and invoking various compute and memory optimizations. When combined, many of these strategies have complex interactions regarding the final training efficiency. Prior work tackling this problem did not have access to the latest set of optimizations, such as FlashAttention or sequence parallelism. In this work, we conduct a comprehensive ablation study of possible training configurations for large language models. We distill this large study into several key recommendations for the most efficient training. For instance, we find that using a micro-batch size of 1 usually enables the most efficient training layouts. Larger micro-batch sizes necessitate activation checkpointing or higher degrees of model parallelism and also lead to larger pipeline bubbles. Our most efficient configurations enable us to achieve state-of-the-art training efficiency results over a range of model sizes, most notably a Model FLOPs utilization of 70.5% when training a Llama 13B model.

1 Introduction

The number of parameters and computational resources spent on training deep neural networks is growing rapidly [1, 3, 14]. The largest models consisting of hundreds of billions of parameters do not even fit onto a single hardware accelerator. Thus, training these models requires various ways of reducing the memory requirements, such as ZeRO [16], activation checkpointing [2], and 3D-parallel (data, tensor, and pipeline parallel) training [13]. 3D parallelism, in particular, has been demonstrated to be effective for the training of Transformer-based large language models (LLMs) with hundreds of billions of parameters [13].

However, training these models efficiently with 3D parallelism requires significant domain expertise and extensive manual effort to determine the ideal configurations. These configurations not only need to combine data, model, and pipeline parallelism most efficiently, but also consider complex interactions with other memory and compute optimizations. FlashAttention [5] in particular has had a notable impact since its release, enabling us to train models at previously impossible degrees of training efficiency. In light of these developments, we conduct a systematic study via a large-scale training efficiency sweep of these interactions. We consider up to 256 GPUs and Llama models with up to 65 billion parameters.

Workshop on Advancing Neural Network Training at 37th Conference on Neural Information Processing Systems (WANT@NeurIPS 2023).

Gefördert durch:

Bundesministerium für Forschung, Technologie und Raumfahrt

Zugangsanfrage
aisc.hpi.de

Infrastruktur

- **Kostenloser Zugang**
- Kein Produktionsbetrieb
  - Daten sollten **anonymisier** oder **synthetisierd** werden
  - Kein **hosting** von Produkten
- **Berichterstattung** & **Veröffentlichung** durch Nutzer:innen
- **Alte Rechte** verbleiben bei Nutzer:innen
- **Neue Rechte** verbleiben bei Nutzer:innen
  - Gewährung von Nutzungsrechten für Forschung und Lehre

Training
- 64 NVIDIA H100 GPU

Inferenz
- 40 NVIDIA A30 GPU

ARM Server
- Ampere Altra Max M128-30 CPU
- 2 x NVIDIA L40 GPUs

GPU Server
- AMD Epyc CPU
- 8 x NVIDIA L40S GPU

Edge
- ARMv8 CPU
- NVIDIA Jetson AGX Module

Neuromorph
- 288 SpiNNaker2 Chips

Storage
- 1.5 PB NVMe

Network
- 400 Gb/s Infiniband
- 200 Gb/s Ethernet

Gefördert durch:
Bundesministerium für Forschung, Technologie und Raumfahrt

Newsletter

Bildung

Talks

tele-task.de/series/1463

- Gastvorträge über Forschung und Innovation

Work shops

aimaker.community

- Praktische Themen
- Beispielthemen: Speech2summary, Docker für ML, Semantische Suche

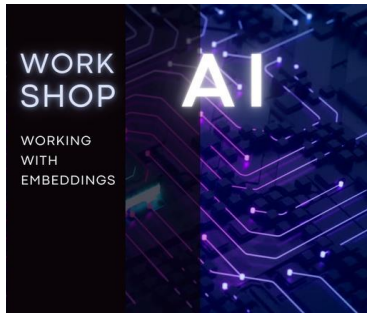

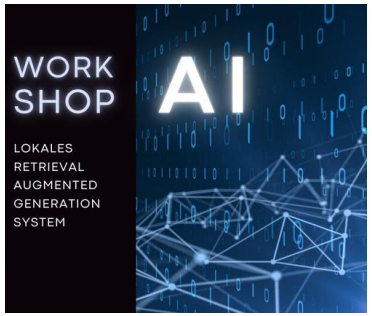

MOOCs

open.hpi.de/channels/ai-service-center

- ChatGPT: Was bedeutet generative KI für die Gesellschaft?
- Profitable AI
- KI Biases verstehen und vermeiden

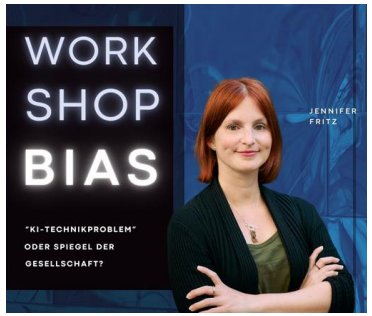

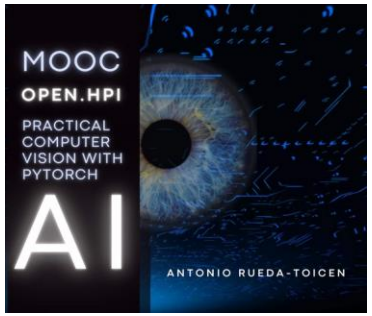

Gefördert durch:
Bundesministerium für Forschung, Technologie und Raumfahrt

Sprechstunde buchen

Consulting

AI consultation hours

- Answering questions:
  - about AI infrastructure
  - about AI models & frameworks
  - about AI use cases

AI pilot projects

- Co-development of a prototype
- Application every three months
- Selection criteria, e.g. AI maturity, public benefit
- Publication of results

Cooperation

- Jointly organized network meetings

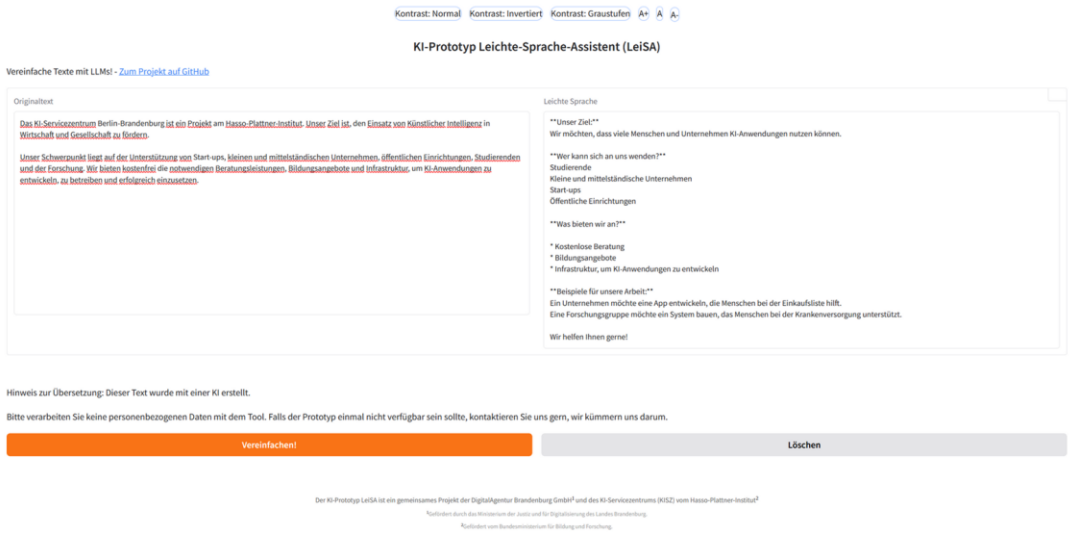

github.com/aihpi/leichte-sprache

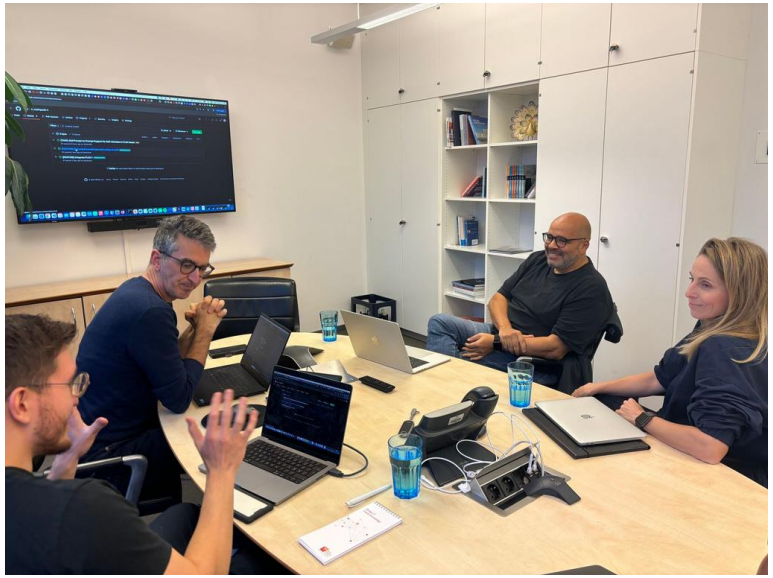

Previous AI pilot projects

- Generation of math problems
- Generation of easy language
- Generation of upcycling suggestions
- Reduction of food waste
- Dating by handwriting

Jetzt bewerben!

Gefördert durch:
Bundesministerium für Forschung, Technologie und Raumfahrt

Verarbeitung gesprochener Sprache

https://tts.aisc.hpi.de/
user: event
pw: aisc@hpi2022!

Prototypen

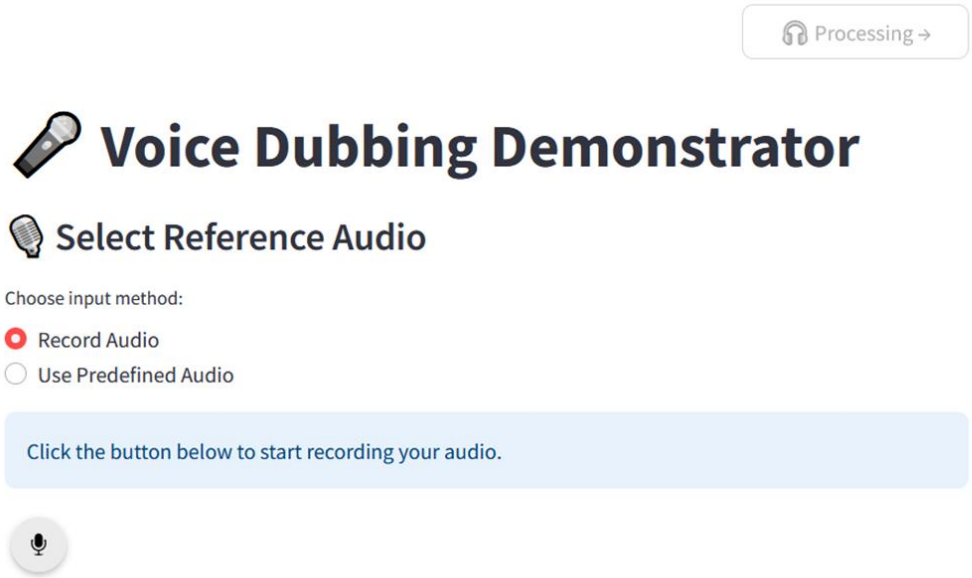

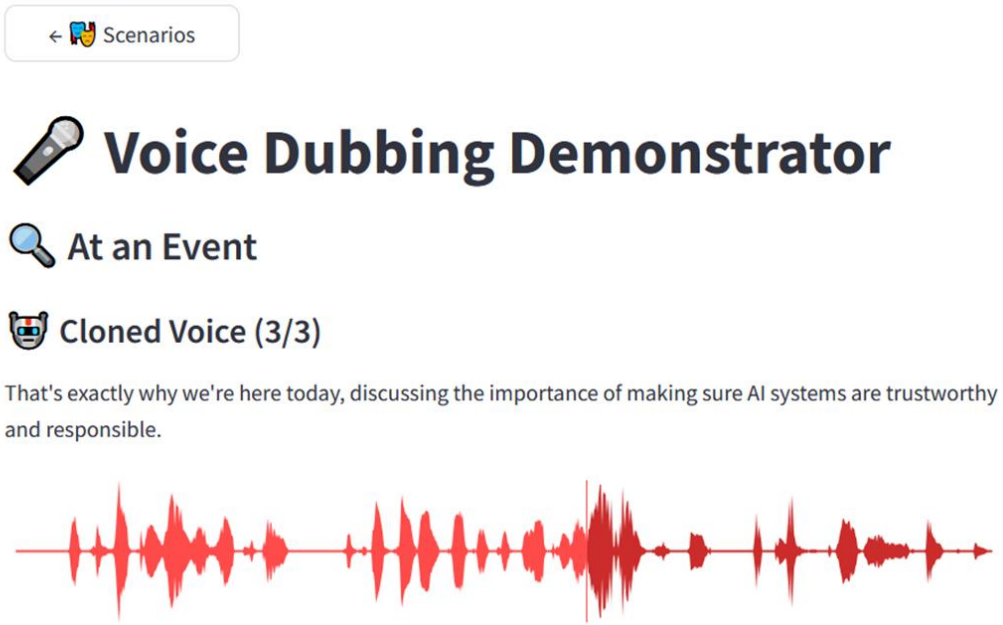

Gefördert durch:
Bundesministerium für Forschung, Technologie und Raumfahrt

KI Service Zentrum
by Hasso-Plattner-Institut

Anwesenheitsliste

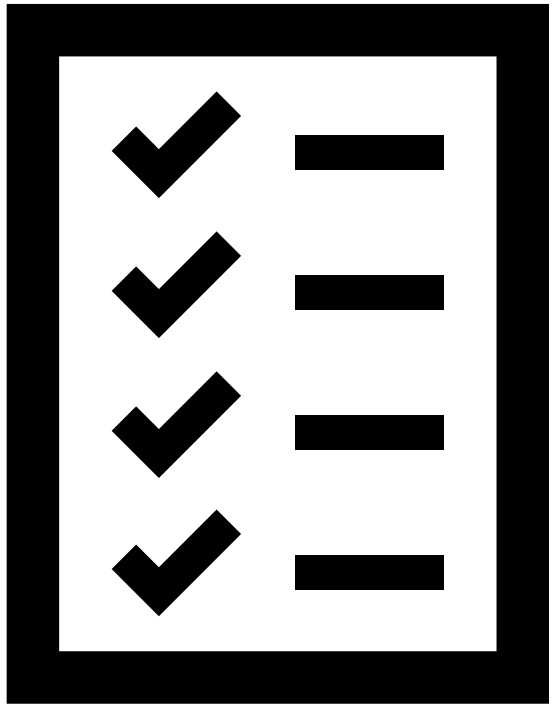

Gefördert durch:

Bundesministerium
für Forschung, Technologie
und Raumfahrt

Nachtrag RAG I – 30 Minuten

Gefördert durch:

Bundesministerium
für Forschung, Technologie
und Raumfahrt

KI Service Zentrum by Hasso-Plattner-Institut HPI

Lernziele

- Was sind *Chunks* und warum muss ich meine Daten ‘chunken’?
  - Zeichenbasiertes Chunking
  - Chunking basierend auf Sinneinheiten
  - Überlappungen
- Welche *Embedding*-Modelle sind für welche Daten geeignet?
  - Text-Embeddings
  - Bild-Embeddings
- Worauf muss ich beim Arbeiten mit Real World Daten achten?
- Verschiedene Dateiformate (z.B. Json, XML) & Tools (z.B. OpenAPI, docling)

Gefördert durch:

Bundesministerium für Forschung, Technologie und Raumfahrt

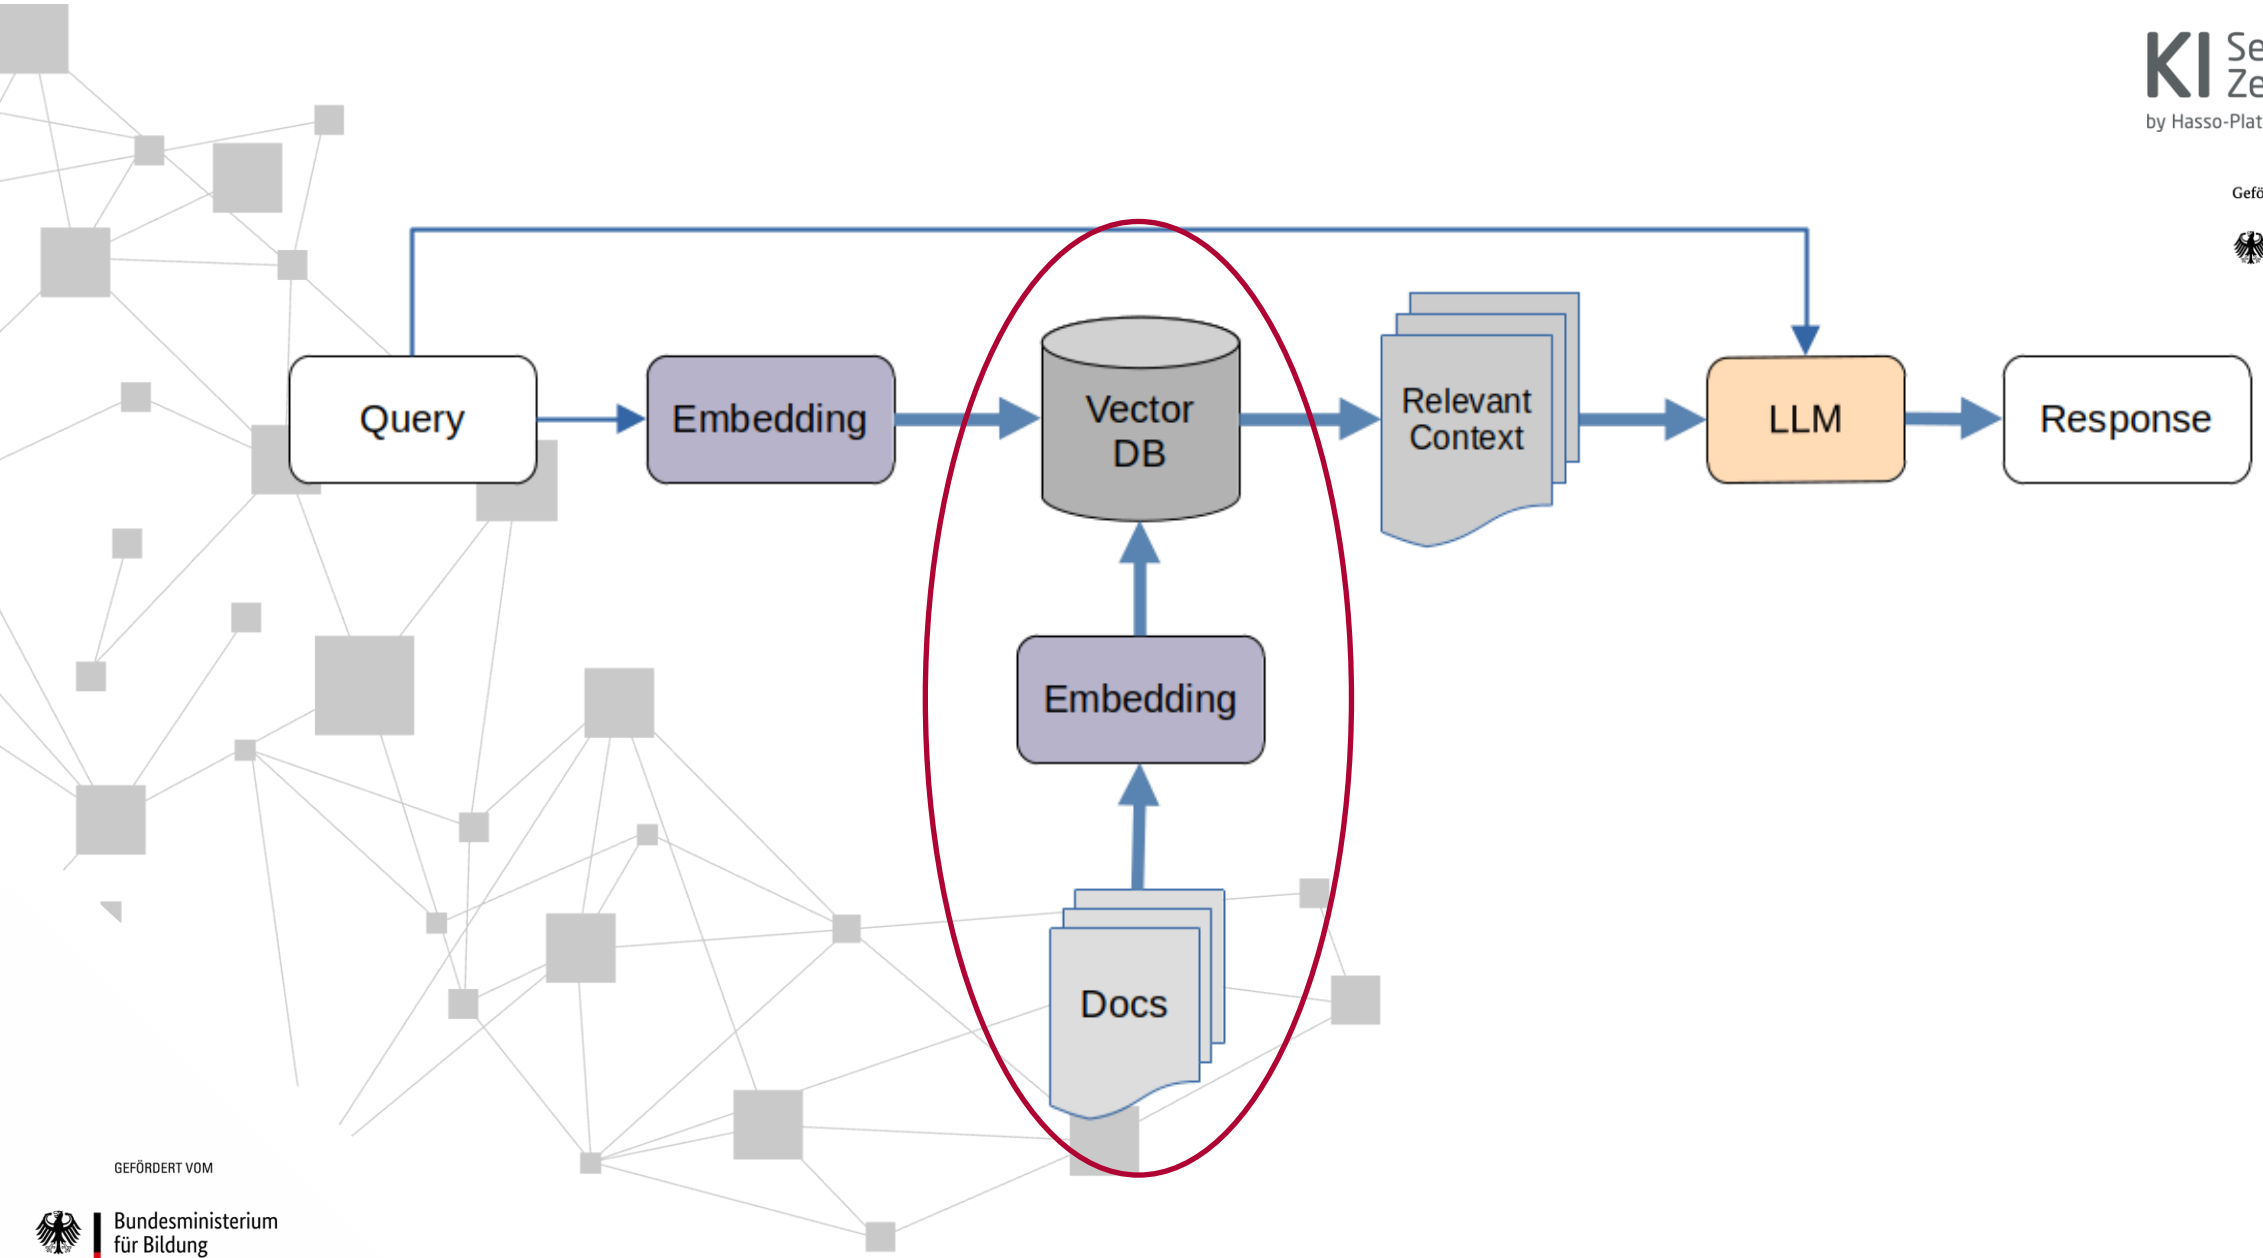

KI Service Zentrum
by Hasso-Plattner-Institut
HPI

Gefördert durch:
Bundesministerium für Forschung, Technologie und Raumfahrt

GEFÖRDERT VOM
Bundesministerium für Bildung und Forschung

Jupyter Notebooks

- Schrittweises Ausführen von Code möglich
  - ‘Interaktives Programmieren’ mit Einblick in Zwischenschritte
  - Einfaches Anpassen/Austauschen von Parameters
- Grafiken werden direkt unter dem Code gerendert
- Code lässt sich mit Markdown anreichern

→ Gut nachvollziehbarer Code; gut geeignet für Demonstrationen

Achtung: Notebooks sind nicht geeignet für (große) Produktionssysteme u. komplexere Software-Projekte

Gefördert durch:
Bundesministerium für Forschung, Technologie und Raumfahrt

Notebook Umgebung

1. Gehe ins Terminal
2. Navigiere zu dem Order, wo das RAG-V2 Repository is, z.B.

```
cd ~/workshops/workshop-ragV2
```

3. Lade den neusten Code

```
git pull
```

4. Navigiere in den Ordner `notebooks`

```
cd ~/workshops/workshop-ragV2/notebooks
```

5. Installiere die virtuelle Umgebung

```
uv venv; source .venv/bin/activate; uv sync
```

Gefördert durch:

Bundesministerium
für Forschung, Technologie
und Raumfahrt

Notebooks

1. Öffne VS Code
2. Öffne das gewünschte Notebook (File-Browser links)
3. Wähle den Kernel aus .venv aus: “rag-notebooks”
4. Klicke die erste Zelle an und drücke str+Enter / cmd+Enter

Gefördert durch:

Bundesministerium
für Forschung, Technologie
und Raumfahrt

00 - Setup

This notebook prepares the local workshop environment:

- Create and configure `.env` with your API key
- Install/sync dependencies with `uv`
- Start Qdrant in Docker and verify it is reachable

In [ ]:

```
from pathlib import Path

def find_repo_root(start: Path) -> Path:
    current = start.resolve()
    for candidate in [current, *current.parents]:
        if (candidate / 'docker-compose.yml').exists():
            return candidate
    raise RuntimeError('Could not find repository root (docker-compose.yml missing).')

REPO_ROOT = find_repo_root(Path.cwd())
NOTEBOOKS_DIR = REPO_ROOT / 'notebooks'
ENV_PATH = NOTEBOOKS_DIR / '.env'
print('Repo root:', REPO_ROOT)
print('Notebooks dir:', NOTEBOOKS_DIR)
print('.env path:', ENV_PATH)
```

1) Configure `.env` (API key + API base URL)

Set your API credentials in `notebooks/.env`.

Expected keys for notebook experiments:

- `OPENAI_API_KEY`
- `OPENAI_API_BASE` (here: `https://api.aisc.hpi.de/`)

> Do not commit real secrets to git.

In [ ]:

```
from pathlib import Path

template = '''# Notebook environment
OPENAI_API_KEY=your_openai_api_key_here
OPENAI_API_BASE=https://api.aisc.hpi.de/
```

RAG Backend Alternative

- Diverse leistungsstarke Sprach- und Embeddingmodelle über KI-Servicezentrum Rechenzentrum verfügbar, z.B.
  - gpt-oss-120b (großes LLM)
  - octen-embedding-8b (leistungsstarkes embedding Modell)
  - Qwen3-vl-embedding-8 (visual language embedding Modell)
- Zugang über API-Key; erfordert wiederum Zugang zum Cluster
- Nutzungsrechte können angefragt werden für tiefere Auseinandersetzung mit RAG (z.B. im Rahmen von Forschung)

Gefördert durch:

Bundesministerium
für Forschung, Technologie
und Raumfahrt

KI Service Zentrum
by Hasso-Plattner-Institut
HPI

kisz@hpi.de

hpi.de/kisz

Deine Meinung ist wichtig!

QR code zum Feedback Formular

Gefördert durch:

Bundesministerium
für Forschung, Technologie
und Raumfahrt